

Project Report

Architecture

Actor

- Fully connected layer - state size to 400
- Fully connected layer - 400 to 300
- Fully connected layer - 64 to 32
- Fully connected layer - 300 to action size

Critic

- Fully connected layer - state size to 400
- Fully connected layer - 400+action size(since we concat) to 300
- Fully connected layer - 64 to 32
- Fully connected layer - 300 to 1

The following parameters were used:

buffer size: 1e6

- batch: 128
- gamma: 0.99
- tau: 1e-3
- learning rate for actor: 1e-4
- learning rate for critic: 1e-4
- weight decay: 0


Learning algorithm

The learning algorithm which was implemented in this project is called [Deep Deterministic Policy Gradient \(DDPG\)](#) in its vanilla form.

Deep Deterministic Policy Gradient (DDPG) is an algorithm which concurrently learns a Q-function and a policy. It uses off-policy data and the Bellman equation to learn the Q-function, and uses the Q-function to learn the policy.

Results

Episode 235 Average Score: 30.04 Environment solved in 135 episodes! Average Score: 30.04



Future ideas

We could try to use the following algorithms next time:

- PPO
- A3C
- D4PG